

CXone Feedback Intelligence

Feedback Management, Reimagined for the AI Era

The contact center world has changed. The data we capture on every interaction is richer than ever, and customers now expect every channel to feel intelligent. That's the opening for a new generation of Feedback Management:

- AI-first by design, not bolted on
- Starts a conversation, not a form
- Context-aware, using what just happened on the interaction
- One smart question at the right moment, instead of a generic survey
- Designed for both AI and human agents, embedded directly within each interaction channel to maximize response rates

The promise sounds simple. Making it real is where the intelligence lives.

The Intelligence Behind the Question

The hard part isn't asking a question. The hard part is knowing which question is worth asking.

Every interaction generates a rich signal: the transcript, the channel, the journey type, the sentiment shift across the call, the topics that came up, whether the customer was frustrated or satisfied, whether the issue was resolved, and whether the agent was human or AI. The new Feedback Management reads all of that in real time and identifies what actually matters on this specific interaction:

- **Topic extraction** that pulls the real subjects of the conversation, not a generic category
- **Sentiment analysis** based on the customer's problem severity, not just whether the agent was polite
- **Journey awareness** that knows if the customer was bounced between bot and human, and how that shaped their experience
- **Anomaly detection** that flags when something doesn't add up — like a textbook-perfect call that still left the customer unhappy

From all of that, the system picks the single most relevant thread and turns it into one contextual question the customer can actually answer. Not "rate your experience." Something like: *"The transfer from the bot took longer than usual — were you happy with the agent's explanation?"*

That's the difference between a survey and a conversation.

From Feedback to Action

What you are looking at is a prototype. The data is illustrative, but the logic behind it is real, and it is what we are building.

Take one 3-star rating from one interaction. Watch what the system does with it:


- It asks a **contextual question** instead of a generic survey — and modeling shows response rates can more than double when the question reflects what actually happened on the call. Customers answer when the question is worth answering.
- From that one rating, the system is designed to surface every similar interaction across the operation, **quantify the operational cost** hiding behind the pattern, and **pinpoint the root cause** — whether it is a routing issue, a policy gap, or an agent behavior.

And it does not stop at the diagnosis. It assigns the work: a coaching plan to the supervisor, a system fix to operations. Specific actions, routed to the people who can act on them.

Decoding the Prototype Logic

Feedback Intelligence — End-to-End Logic Flow


PHASE 1 — ONTOLOGY CONFIGURATION (Admin, one-time setup)




Saved config applies globally across all campaigns — not per-campaign


Interaction ends

PHASE 2 — RUNTIME EXECUTION (Automated)



PHASE 3 — CUSTOMER SURVEY & ANALYSIS





The schema below describes the end-to-end intelligence layer that drives the prototype. Three phases work together to take a raw interaction and turn it into a prioritized, actionable feedback signal.

Phase 1 — Ontology Configuration (Admin Setup)

Before any survey runs, an administrator configures the **Ontology Control Center** — a shared baseline that applies across all campaigns. Each topic type is assigned a **Topic Weight** (1–10, reflecting business criticality) and a **Strategic Modifier** (a multiplier for fine-grained tuning). Sentiment multipliers are also set here: Neutral ×1, Mild ×2, High ×4, Extreme ×5.

A live VU Calculator previews resulting priority tiers in real time. Configuration is saved via **Save Strategic Routing**.

⚠ **Open Decision — Sentiment Model:** The specific sentiment model or provider used to derive Intensity is yet to be decided.

⚠ **Open Decision — Topic Source:** Whether Topic AI will be used for automatic topic extraction from transcripts is yet to be decided.

Phase 2 — Runtime Execution (Per Interaction)

Every time an interaction completes, the pipeline fires:

1. **Eligibility check** — Is this interaction eligible? Channel, agent type, interaction length, and campaign scope are validated.
2. **Suppression rules** — Always-on checks: opt-out list, recency window (30 days), already-responded, internal interaction, open complaint or escalation. If any rule matches, the interaction is suppressed — no survey is sent.
3. **Sampling and filtering** — Volume and rate controls applied per campaign configuration.
4. **Topic identification + VU calculation** — Topics are extracted from the interaction. For each topic, the **VU (Validation Urgency) Score** is calculated:

1 $VU = Weight \times Intensity \times Modifier$

Topics are graded A+ (≥50) through D (<20). Only topics above the configured threshold proceed to survey generation.

5. **AI question generation** — The highest-VU topic drives a contextual question. Rating tier determines framing: 1–2 stars → negative validation; 4–5 stars → positive validation; 3 stars → AI questions skipped.
6. **Survey delivery** — Sent via the interaction channel. Real-time alerting fires for high-VU interactions.

Phase 3 — Customer Survey and Analysis

Once the customer responds:

- The **Validation Score** (0–100%) measures how well the customer's answers align with what the AI predicted. Match on expected answer = bonus (×1.25); mismatch = penalty (×0.5). Scores ≥75% indicate high AI confidence.
- **Routing logic** combines VU and Validation Score to determine the action path:

VU	Validation Score	Outcome
≥35	≥75%	Supervisor escalation
≥35	<75%	AI review queue
<35	≥75%	Analytics

<35	<75%	Standard flow
-----	------	---------------

- **Policy protection** — Keywords like "refund policy", "30-day", or "policy" flag interactions for Operations routing instead of agent coaching.
 - The loop closes: coaching plans go to supervisors, system fixes go to operations, and every signal feeds the dashboard for pattern detection across the operation.
-

Sub-Pages

Detailed technical documentation for each component lives in the sub-pages below:

- [VU Score - Validation Urgency](#) — Formula, components, thresholds, examples, and admin configuration
- [Topic Extraction](#) — Primary and secondary extraction methods, schema constraints, pipeline phases
- [Contextual Questions](#) — Question generation logic, A/B structure, validation scoring, and the INT-022 voice emulator